

A
Capstone Project Synopsis
On

**Pneumonia Disease Detection Using Machine Learning:
RespiraCheck**

Submitted during 4th semester in partial fulfilment of the requirements
for the award of degree of

Bachelor of Technology

in

Computer Engineering

by

Tanishq (24001003129)

Under supervision of

Dr. Umesh



**Department of Computer Engineering
FACULTY OF INFORMATICS & COMPUTING**

**J.C. Bose University of Science & Technology, YMCA
Faridabad – 121006**

Jan-May 2026

TABLE OF CONTENTS

Chapter 1: Introduction

- 1.1 Brief of the Project
- 1.2 Technologies Used
- 1.3 Objectives of the Project

Chapter 2: Proposed Methodology

Chapter 3: Tools and Technologies

Chapter 4: Applications of the Project

- 4.1 Applications
- 4.2 Scope and Limitations

Chapter 5: Conclusion

Chapter 6: Future Work

References

CHAPTER 1: INTRODUCTION

Pneumonia is a serious respiratory infection that inflames the air sacs in one or both lungs, and timely diagnosis is critical for effective treatment. Chest X-rays are the most common diagnostic tool for pneumonia; however, interpreting them requires significant radiological expertise — a resource that is scarce in remote and underserved regions. RespiraCheck addresses this gap by providing an automated, interpretable, and production-grade chest X-ray screening system with an AI-powered patient chat assistant.

1.1 BRIEF OF THE PROJECT

RespiraCheck is a full-stack deep learning application for binary classification of chest X-ray images (NORMAL vs. PNEUMONIA). Two backbone architectures — DenseNet121 and EfficientNetB3 — are trained in parallel using a two-phase transfer learning strategy and automatically evaluated on a held-out test set. The model with the higher AUC-ROC score is deployed; in completed testing, DenseNet121 achieved 88.78% accuracy, AUC-ROC of 0.9549, and 92.31% pneumonia recall, exceeding all stated targets.

Grad-CAM++ heatmaps are generated for every prediction, visually highlighting the lung regions that drove the classification. A Google Gemini 2.5 Flash powered chat assistant is integrated via a /chat API endpoint, enabling patients to ask questions about their result and learn about pneumonia in accessible language. The final deliverable is a Flask REST API (/predict, /chat, /health) paired with a React frontend.

1.2 TECHNOLOGIES USED

- Programming Language: Python 3.10+
- Deep Learning: TensorFlow 2.20.0, tf_keras 2.20.0
- Architectures: DenseNet121, EfficientNetB3 (ImageNet weights)
- Explainability: Grad-CAM++ via tf-keras-vis 0.8.7
- AI Chat: Google Gemini 2.5 Flash (google-genai 1.68.0)
- Data / Visualisation: NumPy, Pandas, Scikit-learn, Matplotlib, Seaborn, OpenCV, Pillow
- Backend / Frontend: Flask 3.1.0, Flask-CORS 5.0.1, React (JavaScript / TypeScript)
- Dev Environment: VS Code, Google Colab, Jupyter Notebook

1.3 OBJECTIVES OF THE PROJECT

- Develop and compare DenseNet121 and EfficientNetB3 for pneumonia classification using a standardised pipeline.
- Implement two-phase transfer learning: frozen-backbone head training followed by selective fine-tuning of the top 30 layers.
- Automate model selection based on AUC-ROC and deploy the best model via Flask REST API and React frontend.
- Integrate Grad-CAM++ so every prediction is accompanied by a visual heatmap overlay.

- Integrate a Gemini 2.5 Flash AI chat assistant for patient education and result explanation.
 - Achieve test accuracy > 80%, pneumonia recall > 93%, and AUC-ROC > 0.92.
-

CHAPTER 2: PROPOSED METHODOLOGY

The pipeline is executed in six structured phases:

1. DATA ACQUISITION AND PREPROCESSING

The Kaggle Chest X-Ray Images (Pneumonia) dataset (5,856 images; 1,583 NORMAL, 4,273 PNEUMONIA) is used. The default 16-image-per-class validation split is discarded; train/ and val/ are merged and re-split 80/20 with stratification. Images are resized to 224×224 pixels, normalised to [0,1], and augmented with random rotations ($\pm 20^\circ$), horizontal flips, zoom (10%), and shifts (10%). Class weights are computed via sklearn to address the 1:2.7 imbalance.

2. MODEL DEVELOPMENT

Both DenseNet121 and EfficientNetB3 are loaded with ImageNet weights and an identical classification head is appended: GlobalAveragePooling2D → Dense(512, ReLU) → Dropout(0.5) → Dense(1, Sigmoid). This shared head ensures any performance difference is attributable solely to the backbone architecture.

3. TRAINING AND FINE-TUNING

Phase 1: The backbone is fully frozen; the classification head is trained for up to 10 epochs (Adam, LR=1e-4, binary cross-entropy). Phase 2: The last 30 layers of the backbone are unfrozen (Batch Normalisation layers remain frozen) and fine-tuned at LR=1e-5. All phases use EarlyStopping (patience=5), ModelCheckpoint (monitoring val_auc), and ReduceLROnPlateau (factor=0.5, patience=3) callbacks.

4. EVALUATION AND AUTOMATIC MODEL SELECTION

Both Phase 2 models are evaluated on the test set using Accuracy, Recall, Precision, F1-Score, AUC-ROC, Average Precision, and Confusion Matrix. ROC and PR curves for both models are plotted on the same axes for comparison. The model with the higher AUC-ROC is automatically written to best_model_path.txt, which the Flask API reads at startup. DenseNet121 was selected: Accuracy 88.78%, AUC-ROC 0.9549, Pneumonia Recall 92.31%, AP 0.9716, F1 91.14%.

5. EXPLAINABILITY WITH GRAD-CAM++

Grad-CAM++ is applied via tf-keras-vis to the final convolutional layer of the winning model, generating a colour heatmap overlay (blue → red: low → high activation) on the original X-ray. The heatmap is returned as a Base64-encoded PNG by the /predict endpoint alongside the prediction label and confidence score.

6. DEPLOYMENT AND AI CHAT INTEGRATION

The Flask REST API exposes: /predict (image upload → prediction + confidence + Grad-CAM heatmap), /chat (Gemini 2.5 Flash multi-turn assistant that explains results, educates on pneumonia, and guides next steps while directing users to consult a doctor), and /health (liveness probe). The React frontend provides the upload interface and chat window. API keys are managed via python-dotenv.

CHAPTER 3: TOOLS AND TECHNOLOGIES TO BE USED

The following table summarises all tools, frameworks, and libraries used in this project:

Category	Technology / Library
Programming Language	Python 3.10+
Deep Learning Framework	TensorFlow 2.20.0, tf_keras 2.20.0
Pre-trained Architectures	DenseNet121, EfficientNetB3 (ImageNet weights)
Explainability	Grad-CAM++ via tf-keras-vis 0.8.7
AI Chat Assistant	Google Gemini 2.5 Flash (google-genai 1.68.0)
Data Handling	NumPy 2.2.4, Pandas 2.2.3, Scikit-learn 1.6.1
Image Processing	OpenCV 4.11.0, Pillow 11.2.1
Visualization	Matplotlib 3.10.1, Seaborn 0.13.2
Backend / Config	Flask 3.1.0, Flask-CORS 5.0.1, python-dotenv 1.2.2
Frontend	React (JavaScript / TypeScript)
Dataset	Kaggle: Chest X-Ray Images (Pneumonia) — 5,856 images
Dev Environment	VS Code, Google Colab, Jupyter Notebook

CHAPTER 4: APPLICATIONS OF THE PROJECT

4.1 APPLICATIONS

- **Clinical Decision Support:** Provides a fast second opinion for radiologists, with Grad-CAM++ heatmaps pinpointing suspicious lung regions and helping prioritise critical cases.
- **Telemedicine and Remote Diagnostics:** Deployable in mobile clinics or rural areas lacking specialist radiologists; only a browser and internet connection are required.
- **Patient Education via AI Chat:** The Gemini 2.5 Flash assistant lets patients ask questions about their diagnosis in plain language, guiding them on next steps within the same interface.
- **Medical Education and Training:** Students and interns can compare their own Xray readings with AI-identified activation regions to reinforce radiographic pattern recognition skills.
- **Scalable Public Health Screening:** Serves as a rapid automated pre-screening layer in high-volume workflows, flagging probable positive cases for priority expert review.

4.2 SCOPE AND LIMITATIONS

Scope: RespiraCheck is a proof-of-concept for binary classification (Pneumonia vs. Normal) on the Kaggle dataset, demonstrating a rigorous, interpretable pipeline that significantly improves on the Semester 3 ResNet50 prototype (74.84% accuracy).

- The system is a pre-diagnostic screening tool and must not replace a qualified radiologist; all predictions require medical verification.
 - Performance may degrade on X-rays from imaging equipment not represented in the training data.
 - Binary classification only — cannot diagnose tuberculosis, COVID-19, pleural effusion, or other thoracic conditions.
 - The /chat endpoint requires an active internet connection and a valid Gemini API key.
-

CHAPTER 5: CONCLUSION

RespiraCheck represents a substantial advancement over the Semester 3 ResNet50 prototype, evolving into a rigorous, production-oriented medical AI system. By training and comparing DenseNet121 and EfficientNetB3 with two-phase transfer learning, automatically selecting the best model, handling class imbalance, incorporating GradCAM++ explainability, and integrating a Gemini 2.5 Flash AI chat assistant, the project addresses the key shortcomings of typical AI medical imaging tools: opacity, singlemodel bias, and lack of patient-facing communication.

The completed evaluation confirms DenseNet121 as the winner with 88.78% test accuracy, AUC-ROC of 0.9549, and 92.31% pneumonia recall — surpassing all stated performance targets. The Flask REST API and React frontend make the system accessible to non-technical medical personnel. While limited to binary pneumonia classification on the Kaggle dataset, the architecture is designed for extensibility to multi-class thoracic disease classification in future iterations.

Note: RespiraCheck is a pre-diagnostic decision support tool and is not intended to replace a qualified radiologist. All predictions must be verified by a medical professional before any clinical action is taken.

CHAPTER 6: FUTURE WORK

- Multi-Class Classification: Extend the model to distinguish bacterial pneumonia, viral pneumonia, COVID-19, tuberculosis, and pleural effusion using the NIH ChestX-ray14 dataset.
- Pneumonia Localisation: Train a YOLOv8 detection model on RSNA bounding box annotations to precisely localise infected lung regions beyond heatmap-level explainability.

- Multimodal AI Chat: Upgrade the Gemini assistant to directly analyse the GradCAM++ heatmap image for image-aware, context-specific explanations.
 - Clinical Validation and Regulatory Pathway: Partner with hospitals for multicentre validation and explore CE/FDA regulatory pathways for clinical deployment.
 - Mobile Application: Build a lightweight mobile version using TensorFlow Lite model quantisation for use on low-resource devices in remote health centres.
 - Federated Learning: Train collaboratively across hospital networks without centralising patient records, addressing HIPAA/GDPR privacy requirements.
-

REFERENCES

1. TensorFlow Documentation (2024). <https://www.tensorflow.org/>
2. Flask Documentation (2024). <https://flask.palletsprojects.com/>
3. Kaggle Chest X-Ray Images (Pneumonia) Dataset (2018).
<https://www.kaggle.com/paultimothymooney/chest-xray-pneumonia>
4. Huang, G., Liu, Z., van der Maaten, L., & Weinberger, K.Q. (2017). Densely Connected Convolutional Networks. CVPR.
5. Tan, M., & Le, Q. (2019). EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks. ICML.
6. Rajpurkar, P. et al. (2017). CheXNet: Radiologist-Level Pneumonia Detection on Chest X-Rays. Stanford ML Group.
7. Selvaraju, R.R. et al. (2017). Grad-CAM: Visual Explanations from Deep Networks via Gradient-Based Localization. ICCV.
8. Chattopadhyay, A. et al. (2018). Grad-CAM++: Generalized Gradient-Based Visual Explanations. WACV.
9. RSNA Pneumonia Detection Challenge (2018).
<https://www.kaggle.com/c/rsnapneumonia-detection-challenge>
10. World Health Organization (2022). Pneumonia Fact Sheet.
<https://www.who.int/news-room/fact-sheets/detail/pneumonia>
11. Scikit-learn Documentation. <https://scikit-learn.org/stable/>
12. tf-keras-vis Library. <https://github.com/keisen/tf-keras-vis>
13. Google Generative AI SDK (google-genai). <https://ai.google.dev/>